

Assignment 3

1. Initilisation and Intial Checks

```
In [1]: #Step 1: Import the Pandas and NumPy Library

import pandas as pd
import numpy as np
# Import the Matplotlib and the Seaborn library
import matplotlib.pyplot as plt
import seaborn as sns

print("Libraries imported successfully!")
```

Libraries imported successfully!

```
In [2]: #Step 2: Mount Google Drive to give Collab access to the dataset

from google.colab import drive

try:
    drive.mount('/content/drive')
    print("Google Drive mounted successfully!")

except Exception as e:
    print(f"An error occurred during mounting: {e}") #Easier to understand
    the error
```

Mounted at /content/drive
Google Drive mounted successfully!

```
In [3]: # Step 3: Read the Kaggle dataset (csv) using read_csv function in Pandas

dataset_path = '/content/drive/MyDrive/Python
Class/Students_Grading_Dataset.csv' #Path from drive set

try:
    df = pd.read_csv(dataset_path, header=0) #Reading the CSV File
    print("CSV file loaded successfully!")

except Exception as e:
    print(f"An unexpected error occurred: {e}") #Easier to understand
the error
```

CSV file loaded successfully!

```
In [4]: # Step 4: Prove it works by displaying top 5
try:
    print("Top 5 rows:\n")
    print(df.head())

    print("\nBottom 5 rows:\n")
    print(df.tail())

except Exception as e:
    print(f"An error occurred while displaying the data: {e}") #Easier to
understand the error
```

Top 5 rows:

	Student_ID	First_Name	Last_Name	Email	Gender
Age \ 0	S1000	Omar	Williams	student0@university.com	Female
22					
1	S1001	Maria	Brown	student1@university.com	Male
18					
2	S1002	Ahmed	Jones	student2@university.com	Male
24					
3	S1003	Omar	Williams	student3@university.com	Female
24					
4	S1004	John	Smith	student4@university.com	Female
23					

	Family_Income_Level	Stress_Level (1-10)	Sleep_Hours_per_Night
0	Medium	5	4.7
1	Medium	4	9.0
2	Low	6	6.2
3	High	3	6.7
4	High	2	7.1

[5 rows x 23 columns]

Bottom 5 rows:

Gender	Age	Student_ID	First_Name	Last_Name	Email
Male	19	4995	S5995	Ahmed	Jones student4995@university.com
Male	19	4996	S5996	Emma	Brown student4996@university.com
Female	24	4997	S5997	John	Brown student4997@university.com
Male	23	4998	S5998	Sara	Davis student4998@university.com
Female	21	4999	S5999	Maria	Brown student4999@university.com

	Department	Attendance (%)	Midterm_Score	Final_Score	...	\
4995	Business	Nan	82.15	60.33	...	
4996	Business	65.11	86.31	49.80	...	
4997	CS	87.54	63.55	64.21	...	
4998	CS	92.56	79.79	94.28	...	
4999	Engineering	83.92	83.24	53.47	...	

	Projects_Score	Total_Score	Grade	Study_Hours_per_Week	\
4995	58.42	85.21	D	25.5	
4996	60.87	95.96	C	5.0	
4997	82.65	54.25	A	24.8	
4998	94.29	55.84	A	16.1	
4999	69.25	77.86	F	29.2	

	Extracurricular_Activities	Internet_Access_at_Home	\
4995	No	Yes	
4996	No	Yes	
4997	Yes	No	
4998	Yes	Yes	
4999	No	Yes	

	Parent_Education_Level	Family_Income_Level	Stress_Level (1-10)	\
4995	High School	Low	10	
4996	NaN	Medium	4	
4997	High School	Medium	4	
4998	Bachelor's	Low	1	
4999	PhD	Low	2	

```
Sleep_Hours_per_Night  
4995           8.3  
4996           4.0  
4997           6.3  
4998           8.4  
4999           6.1
```

[5 rows x 23 columns]

2. Initial Data Exploration

```
In [5]: # Step 5: Find out how many students are in the dataset and how many  
variables are present in the dataset  
  
try:  
  
    print("The row and column count details about the dataset are as  
follows:", df.shape)  
  
    print("\nThe number of student records available are:", df.shape[0])  
#prints first part of the shape function output array  
  
    print("The number of variables are:", df.shape[1]) #prints second part of  
the shape function output array  
  
except Exception as e:  
    print(f"An error occurred while displaying the data: {e}") #Easier to  
understand the error
```

The row and column count details about the dataset are as follows:
(5000, 23)

The number of student records available are: 5000
The number of variables are: 23

```
In [6]: #Step 6: Understanding the columns

try:
    print(df.columns.values)

except Exception as e:
    print(f"An error occurred while displaying the data: {e}") #Easier to
understand the error
```

```
['Student_ID' 'First_Name' 'Last_Name' 'Email' 'Gender' 'Age'
'Department'
'Attendance (%)' 'Midterm_Score' 'Final_Score' 'Assignments_Avg'
'Quizzes_Avg' 'Participation_Score' 'Projects_Score' 'Total_Score'
'Grade' 'Study_Hours_per_Week' 'Extracurricular_Activities'
'Internet_Access_at_Home' 'Parent_Education_Level'
'Family_Income_Level'
'Stress_Level (1-10)' 'Sleep_Hours_per_Night']
```

```
In [7]: # Step 7: Rename variables to more relevant and readable names

try:
    new_column_names = {
        'Student_ID': 'Student ID',
        'First_Name': 'First Name',
        'Last_Name': 'Last Name',
        'Midterm_Score': 'Midterm Score',
        'Final_Score': 'Final Score',
        'Assignments_Avg': 'Assignments Score (Averaged)',
        'Quizzes_Avg': 'Quizzes Score (Averaged)',
        'Participation_Score': 'Participation Score',
        'Projects_Score': 'Projects Score',
        'Total_Score': 'Total Score',
        'Study_Hours_per_Week': 'Study Hours per Week',
        'Extracurricular_Activities': 'Extracurricular Activities',
        'Internet_Access_at_Home': 'Internet Access at Home',
        'Parent_Education_Level': 'Parent Education Level',
        'Family_Income_Level': 'Family Income Level',
        'Stress_Level (1-10)': 'Stress Level (1-10)',
        'Sleep_Hours_per_Night': 'Sleep Hours per Night'
    }
    df = df.rename(columns=new_column_names)
    print("Column names updated successfully!")
    print("The updated Column names are:\n", df.columns.values) #Verify the changes
except Exception as e:
    print(f"An error occurred while renaming columns: {e}")
```

```
Column names updated successfully!
The updated Column names are:
['Student ID' 'First Name' 'Last Name' 'Email' 'Gender' 'Age'
 'Department'
 'Attendance (%)' 'Midterm Score' 'Final Score'
 'Assignments Score (Averaged)' 'Quizzes Score (Averaged)'
 'Participation Score' 'Projects Score' 'Total Score' 'Grade'
 'Study Hours per Week' 'Extracurricular Activities'
 'Internet Access at Home' 'Parent Education Level' 'Family Income
 Level'
 'Stress Level (1-10)' 'Sleep Hours per Night']
```

```
In [8]: # Step 8: Create a Dictionary called Variable Details to explain more about each of the variables

variable_details = {
    'Student ID': 'Unique identifier for each student.',
    'First Name': "Student's first name.",
    'Last Name': "Student's last name.",
    'Email': 'Contact email (can be anonymized).',
    'Gender': 'Male, Female, Other.',
    'Age': 'The age of the student.',
    'Department': "Student's department (e.g., CS, Engineering, Business).",
    'Attendance (%)': 'Attendance percentage (0-100%).',
    'Midterm Score': 'Midterm exam score (out of 100).',
    'Final Score': 'Final exam score (out of 100).',
    'Assignments Score (Averaged)': 'Average score of all assignments (out of 100).',
    'Quizzes Score (Averaged)': 'Average quiz scores (out of 100).',
    'Participation Score': 'Score based on class participation (0-10).',
    'Projects Score': 'Project evaluation score (out of 100).',
    'Total Score': 'Weighted sum of all grades.',
    'Grade': 'Letter grade (A, B, C, D, F).',
    'Study Hours per Week': 'Average study hours per week.',
    'Extracurricular Activities': 'Whether the student participates in extracurriculars (Yes/No).',
    'Internet Access at Home': 'Does the student have access to the internet at home? (Yes/No).',
    'Parent Education Level': 'Highest education level of parents (None, High School, Bachelor\''s, Master\''s, PhD).',
    'Family Income Level': 'Low, Medium, High.',
    'Stress Level (1-10)': 'Self-reported stress level (1: Low, 10: High).',
    'Sleep Hours per Night': 'Average hours of sleep per night.'
} #Creating a dictionary to store the explanation of all the variables
print("Variable Details Dictionary-\n")

# Display the variable details in a user-friendly format
for column, description in variable_details.items(): #.items() function allows us to read the key and value individually
    print(f"{column}: {description}")

https://discuss.python.org/t/printing-the-elements-of-a-dictionary-that-is-contained-in-a-list-how/13568
```

Variable Details Dictionary-

Student ID: Unique identifier for each student.
First Name: Student's first name.
Last Name: Student's last name.
Email: Contact email (can be anonymized).
Gender: Male, Female, Other.
Age: The age of the student.
Department: Student's department (e.g., CS, Engineering, Business).

Attendance (%): Attendance percentage (0-100%).
Midterm Score: Midterm exam score (out of 100).
Final Score: Final exam score (out of 100).
Assignments Score (Averaged): Average score of all assignments (out of 100).
Quizzes Score (Averaged): Average quiz scores (out of 100).
Participation Score: Score based on class participation (0-10).
Projects Score: Project evaluation score (out of 100).
Total Score: Weighted sum of all grades.
Grade: Letter grade (A, B, C, D, F).
Study Hours per Week: Average study hours per week.
Extracurricular Activities: Whether the student participates in extracurriculars (Yes/No).
Internet Access at Home: Does the student have access to the internet at home? (Yes/No).
Parent Education Level: Highest education level of parents (None, High School, Bachelor's, Master's, PhD).
Family Income Level: Low, Medium, High.
Stress Level (1-10): Self-reported stress level (1: Low, 10: High).
Sleep Hours per Night: Average hours of sleep per night.

```
In [9]: # Step 9: Identify the quality issues in the dataset to provide a
comprehensive overview of its integrity and completeness.
try:
    print("Columns with Data missing:")
    print(df.isnull().sum())
    print("\nNumber of duplicates:", df.duplicated().sum())
    print("\nCheck for Columns with incorrect data types:")
    print(df.dtypes)
    print("\nCheck for outliers:")
    plt.boxplot(df['Final Score'])
    plt.title('Final Score Boxplot')
    plt.ylabel('Value')
    plt.xticks(rotation=45) # Rotate x-axis labels for better readability
    plt.show()
    print("\nSimilarly.....Checking for Outliers in other numeric
fields:\n")
    for col in df.select_dtypes(include=["number"]).columns:
        Q1 = np.percentile(df[col], 25)
        Q3 = np.percentile(df[col], 75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        # Identify outliers
        outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]

        # Print the number of outliers for the current column
        print(f"Number of outliers in '{col}': {len(outliers)}")

except Exception as e:
    print(f"An unexpected error occurred: {e}")
```

Columns with Data missing:	
Student ID	0
First Name	0
Last Name	0
Email	0
Gender	0
Age	0
Department	0
Attendance (%)	515
Midterm Score	0
Final Score	0
Assignments Score (Averaged)	503
Quizzes Score (Averaged)	0
Participation Score	0
Projects Score	0
Total Score	0
Grade	0
Study Hours per Week	0
Extracurricular Activities	0
Internet Access at Home	0
Parent Education Level	1800

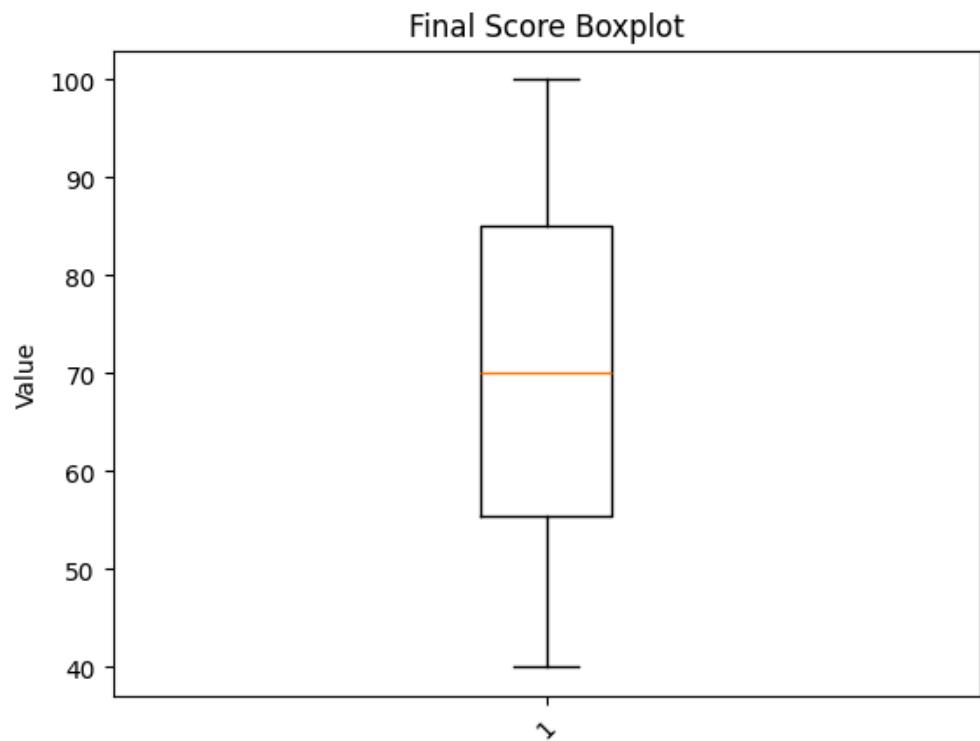
```
Family Income Level          0
Stress Level (1-10)         0
Sleep Hours per Night       0
dtype: int64
```

Number of duplicates: 757

Check for Columns with incorrect data types:

```
Student ID                  object
First Name                  object
Last Name                   object
Email                       object
Gender                      object
Age                          int64
Department                  object
Attendance (%)              float64
Midterm Score               float64
Final Score                 float64
Assignments Score (Averaged) float64
Quizzes Score (Averaged)    float64
Participation Score         float64
Projects Score               float64
Total Score                 float64
Grade                        object
Study Hours per Week        float64
Extracurricular Activities  object
Internet Access at Home    object
Parent Education Level       object
Family Income Level          object
Stress Level (1-10)          int64
Sleep Hours per Night        float64
dtype: object
```

Check for outliers:



Similiarly.....Checking for Outliers in other numeric fields:

```
Number of outliers in 'Age': 0
Number of outliers in 'Attendance (%)': 0
Number of outliers in 'Midterm Score': 0
Number of outliers in 'Final Score': 0
Number of outliers in 'Assignments Score (Averaged)': 0
Number of outliers in 'Quizzes Score (Averaged)': 0
Number of outliers in 'Participation Score': 0
Number of outliers in 'Projects Score': 0
Number of outliers in 'Total Score': 0
Number of outliers in 'Study Hours per Week': 0
Number of outliers in 'Stress Level (1-10)': 0
Number of outliers in 'Sleep Hours per Night': 0
```

```
In [10]: #Step 9.1: Analyze plan of action to handle missing values
try:
    missing_cols = ['Attendance (%)', 'Assignments Score (Averaged)', 'Parent Education Level']
    for col in missing_cols:
        # Calculate missing values within each group and then sum
        missing_counts = df.groupby('Department')[col].apply(lambda x: x.isnull().sum())

        # Calculate total values per department
        total_counts = df.groupby('Department')[col].size()

        # Calculate percentage of missing values
        percentage_missing = (missing_counts / total_counts) * 100

        # Round percentage to two decimal places
        percentage_missing = percentage_missing.round(2)

        print(f"Missing values in {col} by department:\n{missing_counts}\n")
        print(f"Percentage of missing values in {col} by department:\n{percentage_missing}\n")

except Exception as e:
    print(f"An unexpected error occurred: {e}")
```

```
Missing values in Attendance (%) by department:
```

```
Department
Business      111
CS            194
Engineering   152
Mathematics   58
Name: Attendance (%), dtype: int64
```

```
Percentage of missing values in Attendance (%) by department:
```

```
Department
Business      10.69
CS            9.76
Engineering   10.15
Mathematics   12.16
Name: Attendance (%), dtype: float64
```

```
Missing values in Assignments Score (Averaged) by department:
```

```
Department
Business      123
CS            203
Engineering   125
Mathematics   52
Name: Assignments Score (Averaged), dtype: int64
```

```
Percentage of missing values in Assignments Score (Averaged) by department:
```

```
Department
Business      11.85
```

```
CS           10.21
Engineering   8.35
Mathematics  10.90
Name: Assignments Score (Averaged), dtype: float64
```

```
Missing values in Parent Education Level by department:
Department
Business      382
CS            704
Engineering    533
Mathematics   181
Name: Parent Education Level, dtype: int64
```

```
Percentage of missing values in Parent Education Level by department:
Department
Business      36.80
CS            35.41
Engineering    35.60
Mathematics   37.95
Name: Parent Education Level, dtype: float64
```

```
In [11]: # Step 9.1.1: Calculate the overall percentage of missing values in 'Parent Education Level'

missing_parent_education = df['Parent Education Level'].isnull().sum()
total_parent_education = len(df['Parent Education Level'])
percentage_missing_parent_education = (missing_parent_education / total_parent_education) * 100

print(f"Percentage of missing values in 'Parent Education Level': {percentage_missing_parent_education:.2f}%")

# Group by 'Grade' and calculate missing values in 'Parent Education Level'
missing_by_grade = df.groupby('Grade')['Parent Education Level'].apply(lambda x: x.isnull().sum())

# Calculate total students in each grade
total_by_grade = df.groupby('Grade')['Parent Education Level'].count()

# Calculate the percentage of missing values in 'Parent Education Level' for each grade
percentage_missing_by_grade = (missing_by_grade / total_by_grade) * 100

print("\nPercentage of missing 'Parent Education Level' by Grade:")
percentage_missing_by_grade
```

Percentage of missing values in 'Parent Education Level': 36.00%

Percentage of missing 'Parent Education Level' by Grade:

Out[11]:

	Parent Education Level
Grade	
A	57.009346
B	55.694228
C	56.804734
D	52.816901
F	58.733205

dtype: float64

```
In [12]: #Step 9.2: Apply imputation techniques
try:
    for col in ['Attendance (%)', 'Assignments Score (Averaged)', 'Parent
Education Level']:
        if df[col].dtype == 'object':
            mode_val = df[col].mode()[0] # Calculate mode
            df[col].fillna(mode_val, inplace=True) # Impute missing values
with mode
        else:
            mean_val = df[col].mean() # Calculate mean
            df[col].fillna(mean_val, inplace=True) # Impute missing values
with mean

    print("Columns with Data missing:")

    print(df.isnull().sum())

except Exception as e:
    print(f"An unexpected error occurred: {e}")
```

```
Columns with Data missing:
```

```
Student ID          0
First Name          0
Last Name           0
Email               0
Gender              0
Age                 0
Department          0
Attendance (%)      0
Midterm Score       0
Final Score         0
Assignments Score (Averaged) 0
Quizzes Score (Averaged) 0
Participation Score 0
Projects Score       0
Total Score          0
Grade               0
Study Hours per Week 0
Extracurricular Activities 0
Internet Access at Home 0
Parent Education Level 0
Family Income Level 0
Stress Level (1-10)   0
Sleep Hours per Night 0
dtype: int64
```

```
<ipython-input-12-06102eaa7a15>:9: FutureWarning: A value is
trying to be set on a copy of a DataFrame or Series through
chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will
never work because the intermediate object on which we are setting
```

values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(mean_val, inplace=True) # Impute missing values  
with mean  
<ipython-input-12-06102eaa7a15>:6: FutureWarning: A value is  
trying to be set on a copy of a DataFrame or Series through  
chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will  
never work because the intermediate object on which we are setting  
values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(mode_val, inplace=True) # Impute missing values  
with mode
```

In [13]: #step 9.3: Drop the duplicate records

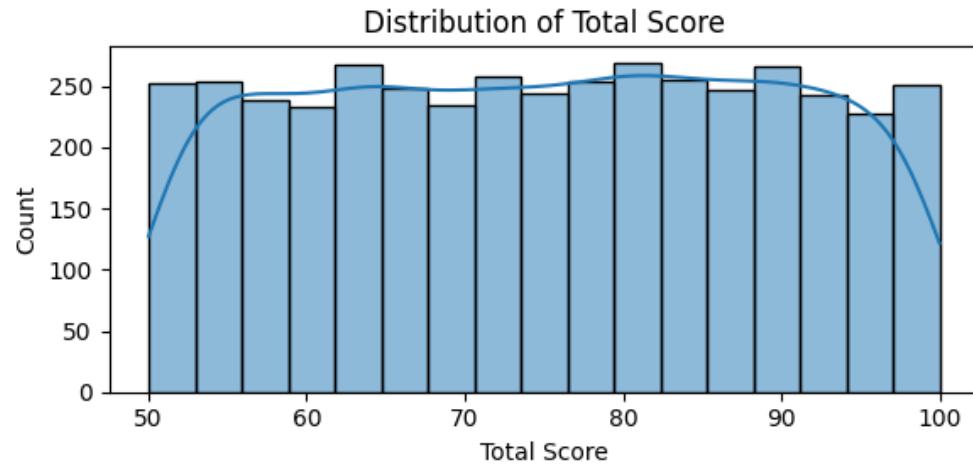
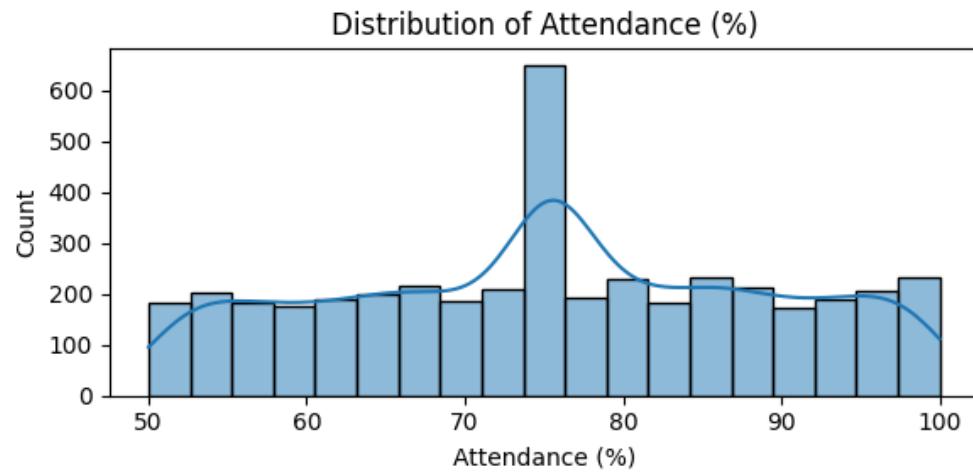
```
try:  
    print("Number of duplicates:", df.duplicated().sum())  
    df.drop_duplicates(inplace=True)  
    print("Number of duplicates after dropping:", df.duplicated().sum())  
    print("Duplicate records removed successfully!")  
    print("Number of rows after dropping duplicates:", df.shape[0])  
except Exception as e:  
    print(f"An unexpected error occurred: {e}")
```

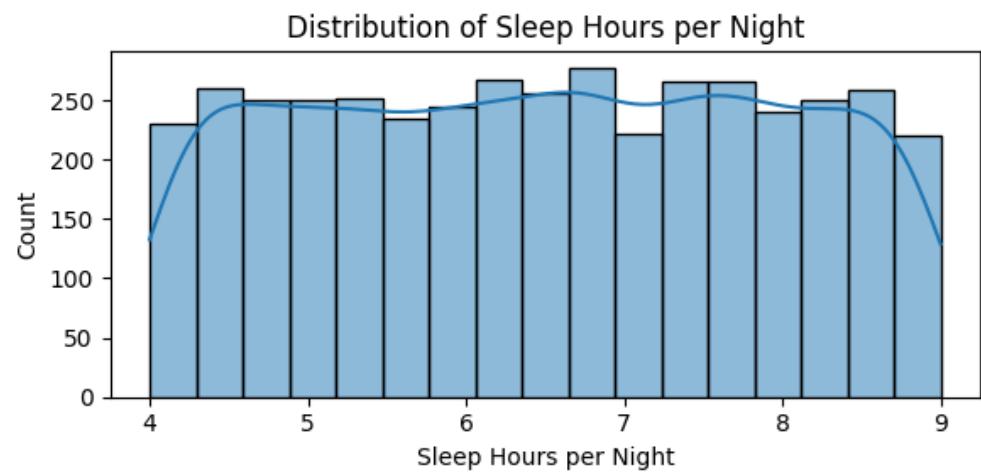
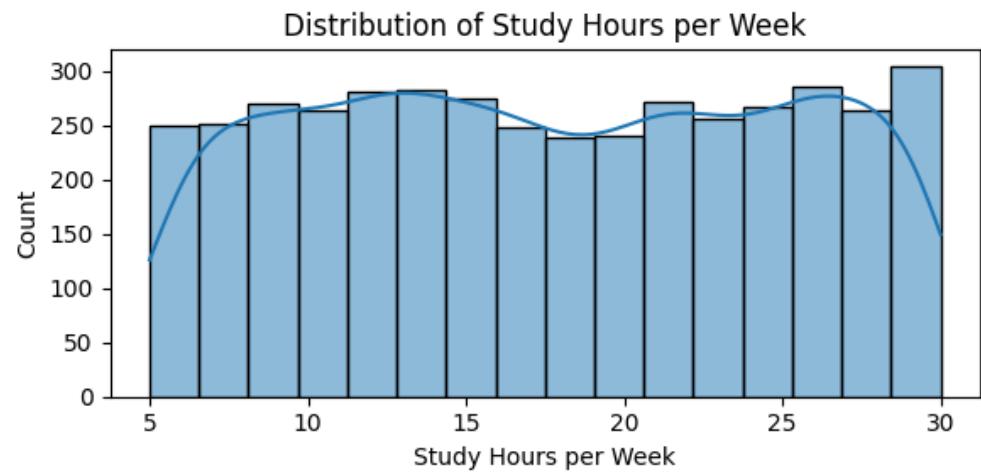
```
Number of duplicates: 757  
Number of duplicates after dropping: 0  
Duplicate records removed successfully!  
Number of rows after dropping duplicates: 4243
```

Visualisation and Analysis

Q.1) What is the distribution of each numerical metric (e.g., attendance, scores, study hours)?

```
In [14]: #What is the distribution of each numerical metric (e.g., attendance, scores, study hours)?  
  
# List of columns to visualize  
numeric_cols = ['Attendance (%)',  
                 'Total Score', 'Study Hours per Week', 'Sleep Hours per Night']  
  
# Loop through each column  
for col in numeric_cols:  
    plt.figure(figsize=(6,3)) # Create a new figure for the plot  
    sns.histplot(df[col], kde=True) # Create a histogram with a density curve  
    plt.title(f'Distribution of {col}') # Add a title to the plot  
    plt.tight_layout() # Adjust the plot layout  
    plt.show() # Show the plot
```





Q.2) How are students distributed across key categories (e.g., gender, department, grade)?

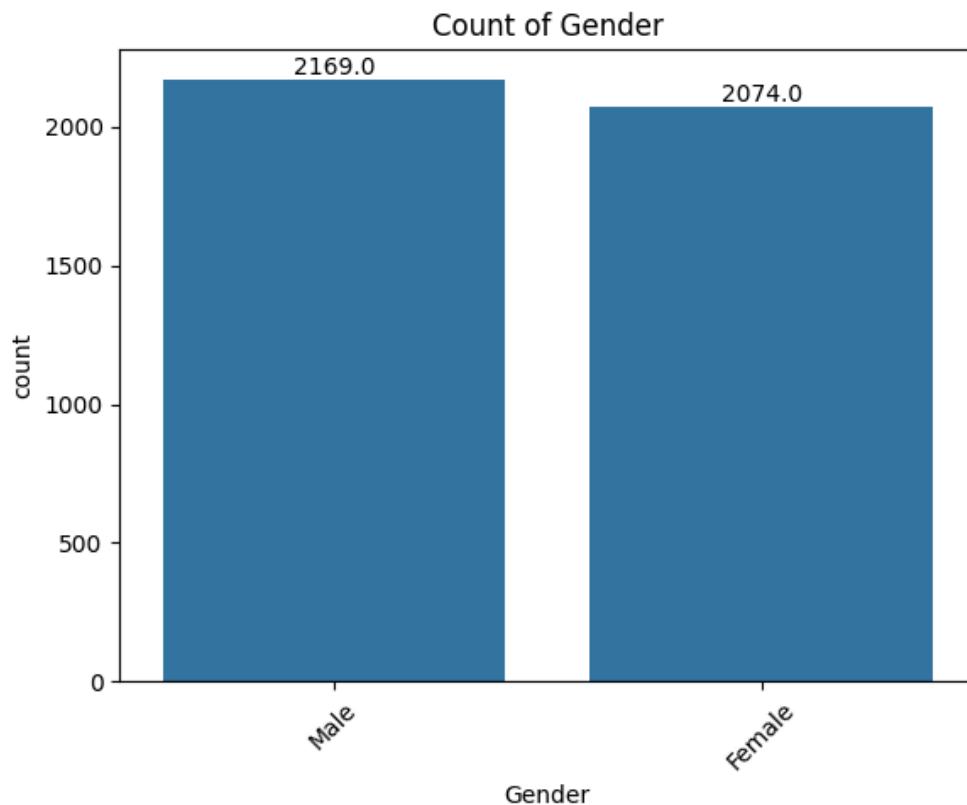
```
In [15]: cat_cols = ['Gender', 'Department', 'Grade',
                 'Extracurricular Activities', 'Family Income Level']

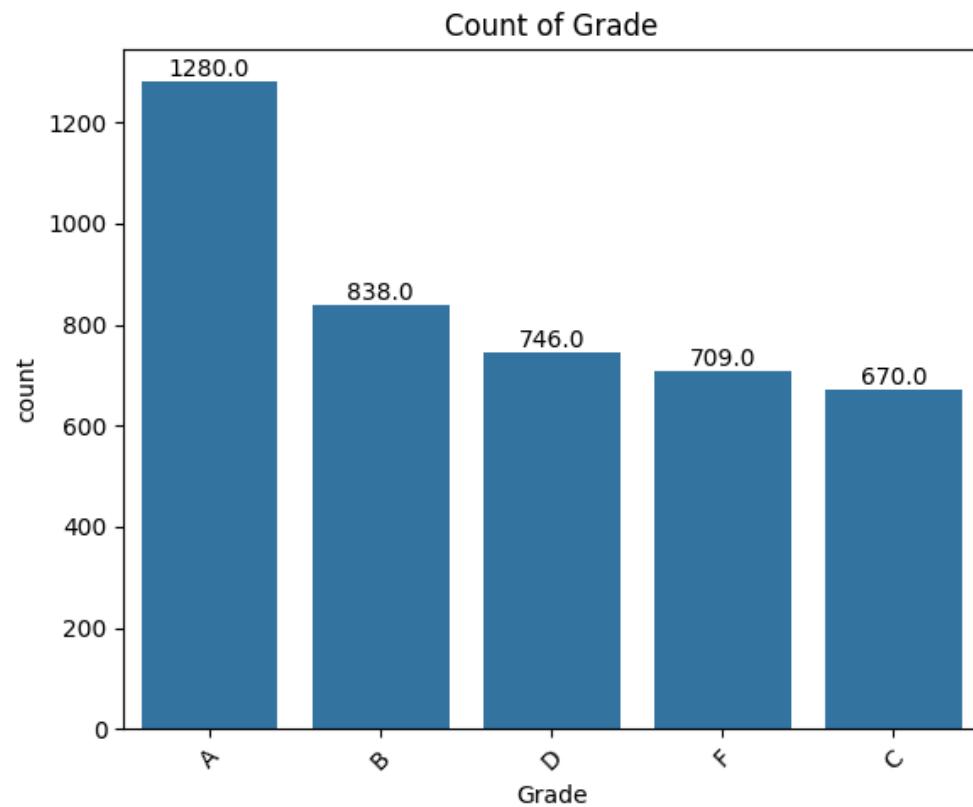
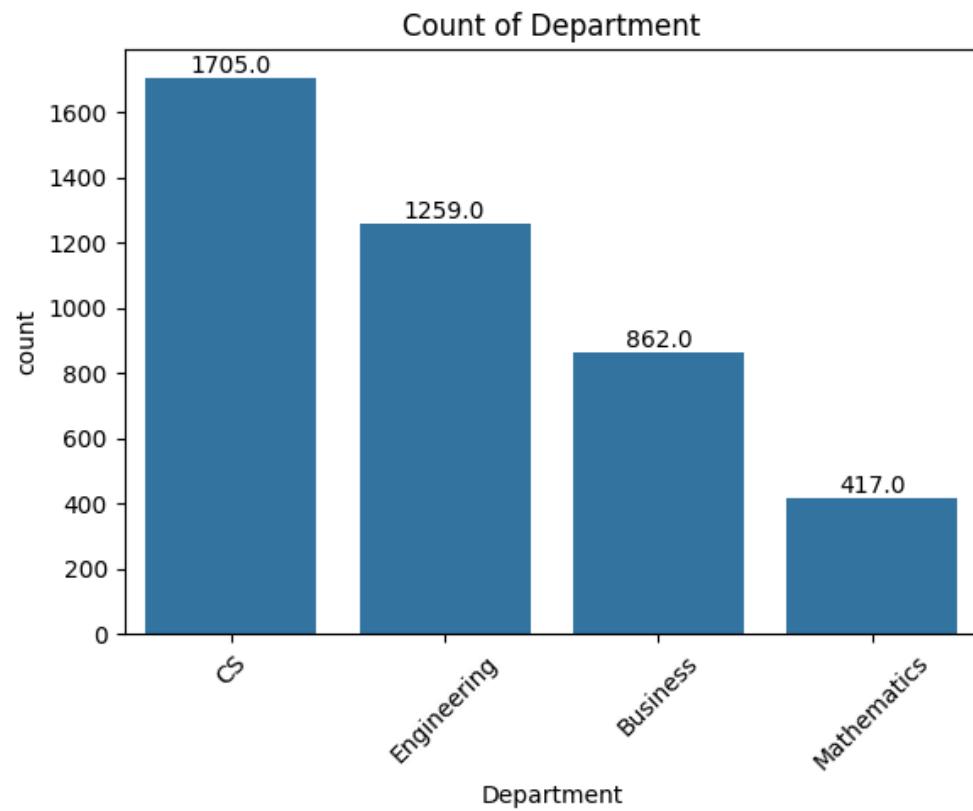
for col in cat_cols:
    plt.figure(figsize=(6,5))
    ax = sns.countplot(data=df, x=col,
                        order=df[col].value_counts().index)
    plt.title(f'Count of {col}')
    plt.xticks(rotation=45)
    plt.tight_layout()

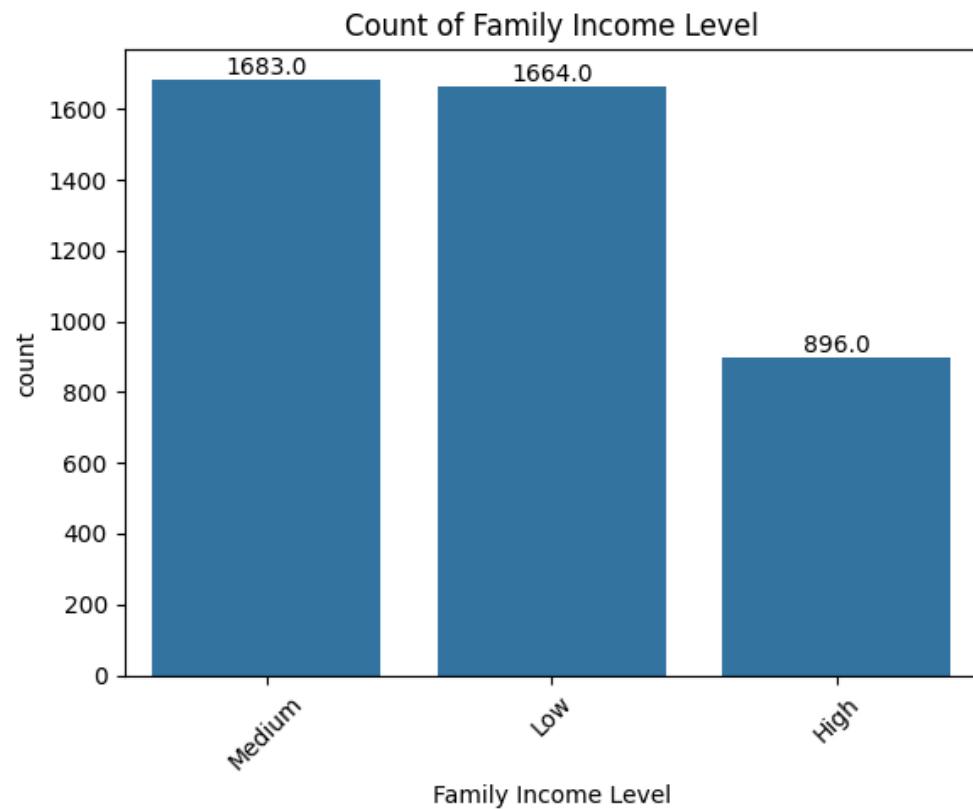
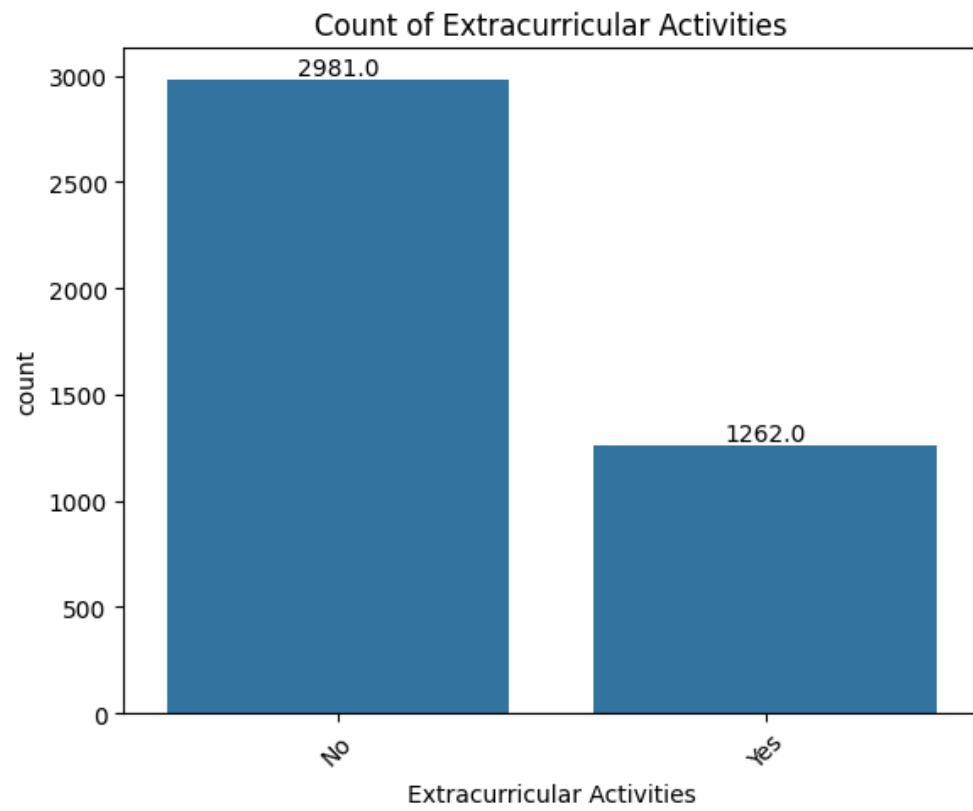
    # Add value counts on top of bars
    for p in ax.patches:
        ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2.,
                                         p.get_height()),
                    ha='center', va='center', xytext=(0, 5),
                    textcoords='offset points')

#https://dataplotplus.com/how-to-annotate-bars-with-values-on-pandas-bar-plots/

plt.show()
```







Q.3) Which department has the highest share of A grades, and which has the largest proportion of D grades?

```
In [16]: # Group data by 'Department' and 'Grade', then count 'Student ID' to get
# student counts per group
grade_counts = df.groupby(['Department', 'Grade'])['Student
ID'].count().unstack(fill_value=0)
# unstack() pivots the 'Grade' level to columns, filling missing values
with 0

# Calculate the total number of students per department (summing across
# grades)
department_totals = grade_counts.sum(axis=1) # axis=1 sums horizontally
(across columns)

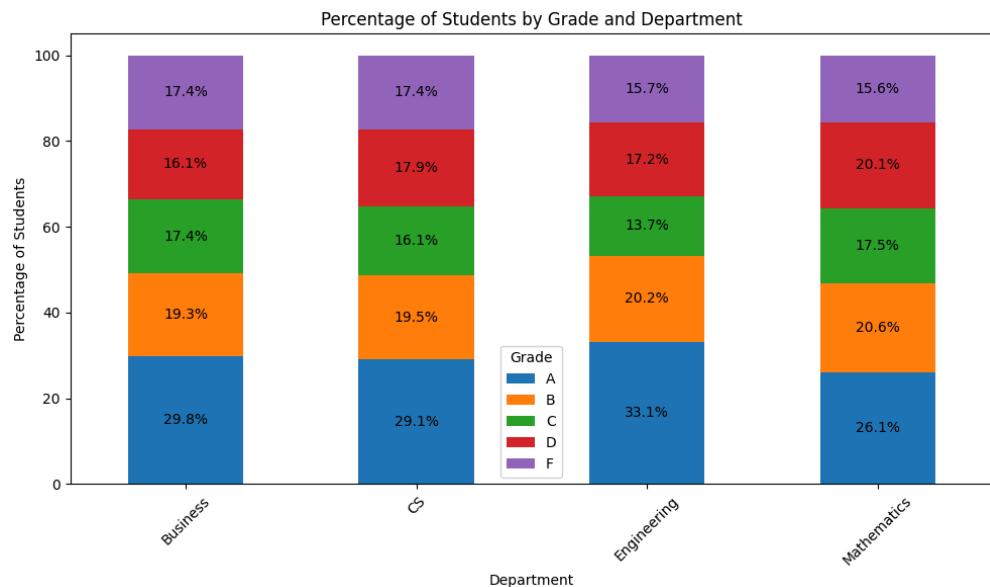
# Calculate the percentage of students with each grade in each department
grade_percentages = grade_counts.div(department_totals, axis=0) * 100 # 
Divides each grade count by department total, then multiplies by 100

# Create the stacked bar chart using the percentage data
ax = grade_percentages.plot(kind='bar', stacked=True, figsize=(10, 6)) #
kind='bar' specifies a bar chart, stacked=True makes it stacked
plt.title('Percentage of Students by Grade and Department')
plt.xlabel('Department')
plt.ylabel('Percentage of Students')
plt.xticks(rotation=45)
plt.legend(title='Grade') # Add a legend with the title 'Grade'

# Annotate the bars with percentages
for p in ax.patches: # Loop through each bar (patch) in the chart
    width, height = p.get_width(), p.get_height() # Get the width and
height of the bar
    x, y = p.get_xy() # Get the x and y coordinates of the bar's bottom
left corner
    if height > 0: # Only annotate if the bar has a height greater than 0
        ax.text(x + width / 2, y + height / 2, f'{height:.1f}%', 
ha='center', va='center') # Add the percentage text to the center of the
bar

#https://dataplotplus.com/how-to-annotate-bars-with-values-on-pandas-bar-
plots/

plt.tight_layout() # Adjust layout to prevent overlapping elements
plt.show() # Display the chart
```

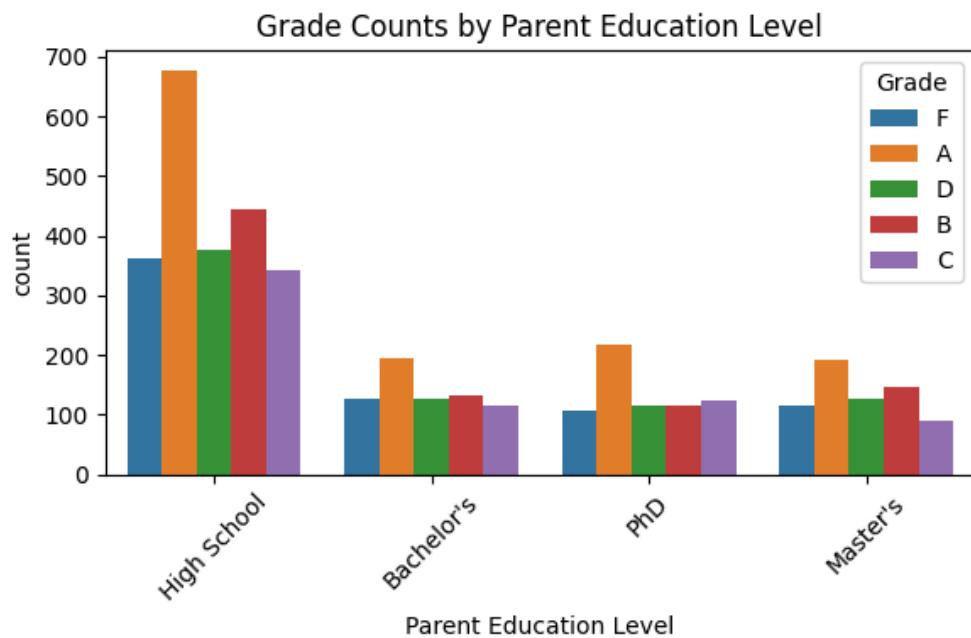


Q.4) How does parental education level relate to student's final letter grades?

```
In [17]: plt.figure(figsize=(6,4))
sns.countplot(
    data=df,
    x='Parent Education Level',
    hue='Grade',
    order=df['Parent Education Level'].value_counts().index # Order
categories by frequency
)
plt.title('Grade Counts by Parent Education Level')
plt.xticks(rotation=45)

#values have not been displayed because they are overlapping

plt.tight_layout()
plt.show()
```



Q.5) Which department has the highest number of high-scoring students who participate in sports?

```
In [18]: # Filter students with Total Score > 80
high_scorers = df[df['Total Score'] > 80]

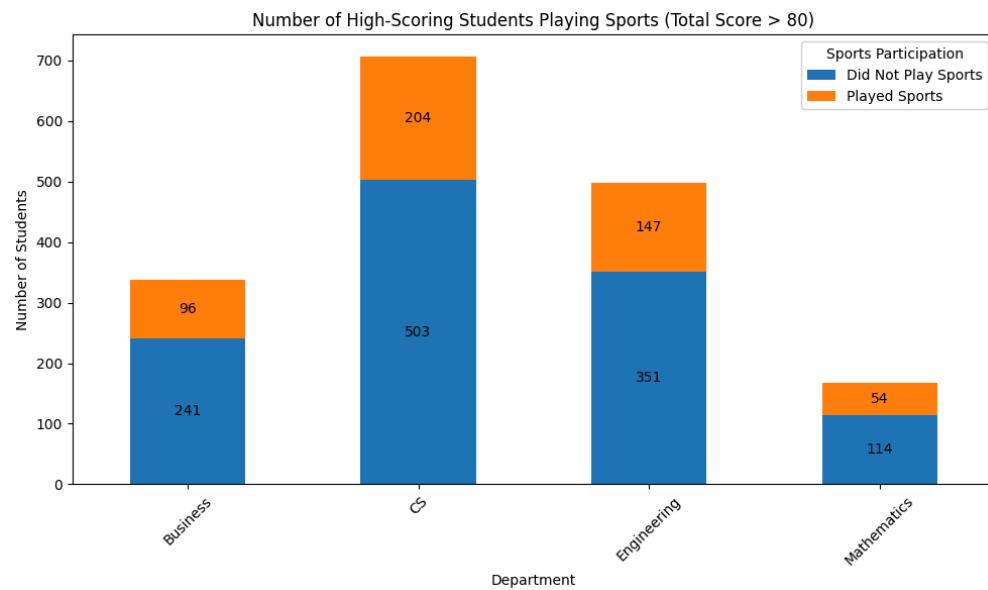
# Group by department and Extracurricular Activities, then count
sports_summary = high_scorers.groupby(['Department', 'Extracurricular Activities'])['Student ID'].count().unstack(fill_value=0)

# Rename columns for clarity
sports_summary = sports_summary.rename(columns={
    'No': 'Did Not Play Sports',
    'Yes': 'Played Sports'
})

# Create the stacked bar chart
ax = sports_summary.plot(kind='bar', stacked=True, figsize=(10, 6))
plt.title('Number of High-Scoring Students Playing Sports (Total Score > 80)')
plt.xlabel('Department')
plt.ylabel('Number of Students')
plt.xticks(rotation=45)
plt.legend(title='Sports Participation', loc='upper right')

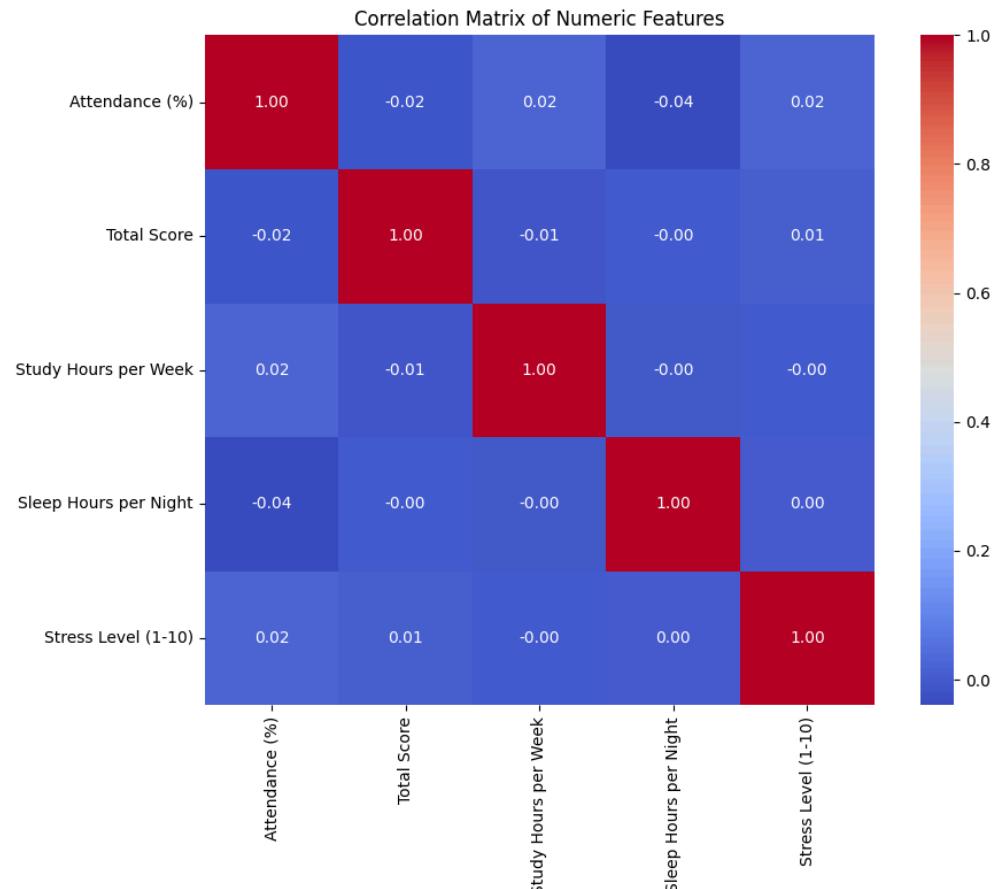
# Display values on the chart
for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.text(x + width / 2, y + height / 2, int(height), ha='center', va='center')

plt.tight_layout()
plt.show()
```



Q.5) What linear relationships exist among the numeric attributes?

```
In [19]: plt.figure(figsize=(10,8))
corr = df[numERIC_COLS + ['Stress Level (1-10)']].corr()
sns.heatmap(
    corr, annot=True, fmt=".2f",
    cmap='coolwarm', square=True
)
plt.title('Correlation Matrix of Numeric Features')
plt.tight_layout()
plt.show()
```



```
In [26]: import statsmodels.api as sm

# Define independent and dependent variables
X = df[['Midterm Score', 'Final Score', 'Assignments Score (Averaged)',
         'Quizzes Score (Averaged)', 'Participation Score', 'Projects
Score',
         'Study Hours per Week', 'Stress Level (1-10)', 'Sleep Hours per
Night',
         'Attendance (%)']] # Add all your desired numeric columns here

y = df['Total Score']

# Add a constant to the independent variables (intercept)
X = sm.add_constant(X)

# Fit the OLS model
model = sm.OLS(y, X).fit()

# Print the regression results
print(model.summary())

#https://www.geeksforgeeks.org/ordinary-least-squares-ols-using-
statsmodels/
```

OLS Regression Results

Dep. Variable:	Total Score	R-squared:
0.002		
Model:	OLS	Adj. R-squared:
-0.000		
Method:	Least Squares	F-statistic:
0.9321		
Date:	Mon, 19 May 2025	Prob (F-statistic):
0.502		
Time:	00:52:50	Log-Likelihood:
-17325.		
No. Observations:	4243	AIC:
3.467e+04		
Df Residuals:	4232	BIC:
3.474e+04		
Df Model:	10	
Covariance Type:	nonrobust	

		coef	std err	t
P> t	[0.025 0.975]			
const		76.8013	2.999	25.613
0.000	70.923	82.680		
Midterm Score		-0.0018	0.013	-0.137
0.891	-0.027	0.024		

Final Score		0.0032	0.013	0.246
0.806	-0.022	0.028		
Assignments Score (Averaged)		0.0048	0.016	0.299
0.765	-0.027	0.037		
Quizzes Score (Averaged)		0.0241	0.015	1.585
0.113	-0.006	0.054		
Participation Score		-0.1257	0.076	-1.643
0.100	-0.276	0.024		
Projects Score		-0.0242	0.015	-1.581
0.114	-0.054	0.006		
Study Hours per Week		-0.0203	0.030	-0.670
0.503	-0.080	0.039		
Stress Level (1-10)		0.0265	0.077	0.343
0.731	-0.125	0.178		
Sleep Hours per Night		-0.0269	0.152	-0.177
0.860	-0.325	0.271		
Attendance (%)		-0.0165	0.016	-1.016
0.309	-0.048	0.015		
<hr/>				
<hr/>				
Omnibus:		3143.530	Durbin-Watson:	
1.995				
Prob(Omnibus):		0.000	Jarque-Bera (JB):	
248.232				
Skew:		-0.013	Prob(JB):	
1.25e-54				
Kurtosis:		1.815	Cond. No.	
2.47e+03				
<hr/>				
<hr/>				

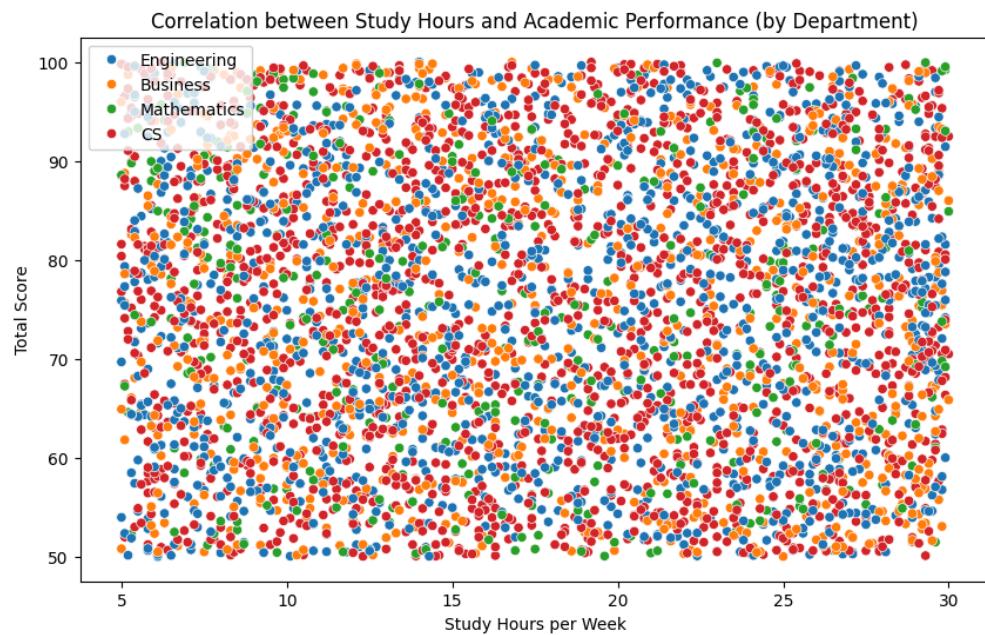
Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.47e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [21]: department_avg_scores = df.groupby('Department')['Total Score'].mean()

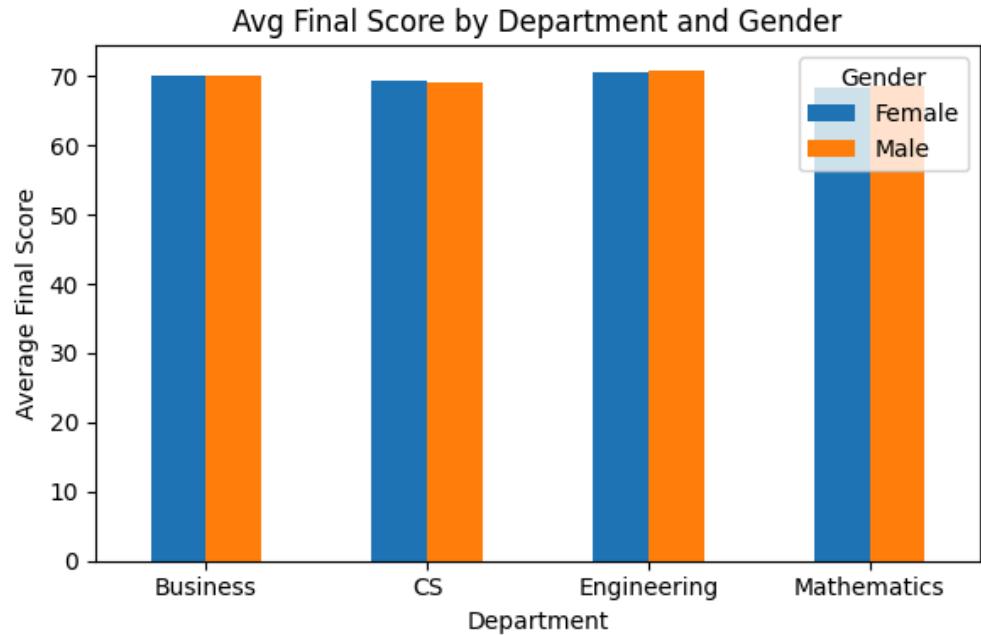
# Create the scatter plot with department as hue and averaged total score
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='Study Hours per Week', y='Total Score',
hue='Department')
plt.title('Correlation between Study Hours and Academic Performance (by
Department)')
plt.xlabel('Study Hours per Week')
plt.ylabel('Total Score')

plt.legend()
plt.show()
```



Q.6) Do average final exam scores differ across departments and between genders?

```
In [22]: grouped = df.groupby(['Department', 'Gender'])['Final Score'].mean().unstack()
grouped.plot(kind='bar', figsize=(6,4))
plt.title('Avg Final Score by Department and Gender')
plt.ylabel('Average Final Score')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```



Q.7) What distinct study and performance patterns can we identify among students based on their academic scores, study hours, and stress levels?

```
In [23]: from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

# Select relevant features for clustering
features_for_clustering = ['Midterm Score', 'Final Score', 'Assignments Score (Averaged)', 'Quizzes Score (Averaged)', 'Study Hours per Week', 'Stress Level (1-10)']
X_cluster = df[features_for_clustering]

# Scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_cluster)

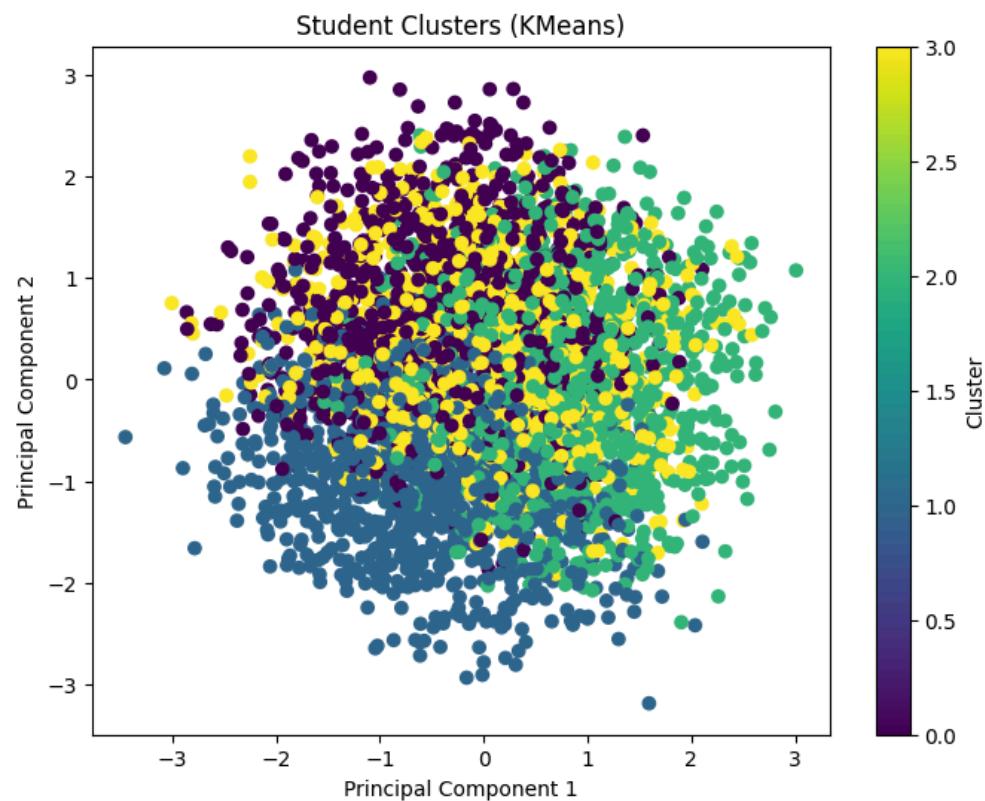
# Apply PCA to reduce dimensionality (optional, but can improve clustering)
pca = PCA(n_components=2) # Reduce to 2 principal components for visualization
X_pca = pca.fit_transform(X_scaled)

# Perform KMeans clustering
kmeans = KMeans(n_clusters=4, random_state=42) # Experiment with different numbers of clusters
df['Cluster'] = kmeans.fit_predict(X_scaled) # Use scaled data for clustering

# Visualize clusters (using PCA-reduced data for 2D visualization)
plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=df['Cluster'], cmap='viridis')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('Student Clusters (KMeans)')
plt.colorbar(label='Cluster')
plt.show()

# Analyze cluster characteristics
for cluster in range(4): # Assuming 4 clusters
    print(f"Cluster {cluster}:")
    print(df[df['Cluster'] == cluster][features_for_clustering].describe())

#https://365datascience.com/tutorials/python-tutorials/pca-k-means/
#https://pulsedatahub.com/blog/k-means-clustering
```

**Cluster 0:**

	Midterm Score	Final Score	Assignments Score (Averaged)	\
count	1059.000000	1059.000000		1059.000000
mean	74.554721	82.324797		75.865203
std	14.824129	10.884766		13.563318
min	40.010000	49.790000		50.080000
25%	63.765000	74.645000		65.545000
50%	76.330000	83.210000		74.828117
75%	86.155000	91.070000		86.390000
max	99.910000	99.980000		99.960000

	Quizzes Score (Averaged) (1-10)	Study Hours per Week	Stress Level
count	1059.000000	1059.000000	
mean	5.556185	71.218593	10.763362
std	2.824984	13.908271	3.629310
min	1.000000	50.030000	5.000000
25%	3.000000	59.390000	7.600000
50%	6.000000	70.100000	10.400000
75%	8.000000	81.790000	13.400000
max	10.000000	99.900000	21.100000

Cluster 1:

	Midterm Score	Final Score	Assignments Score (Averaged) \
count	1058.000000	1058.000000	1058.000000
mean	52.817807	57.252836	74.227787
std	8.512407	10.724053	13.751985
min	40.000000	40.090000	50.130000
25%	45.380000	48.285000	62.817500
50%	52.075000	56.675000	74.828117
75%	59.117500	64.472500	85.102500
max	76.980000	86.790000	99.980000

	Quizzes Score (Averaged)	Study Hours per Week	Stress Level
(1-10)			
count	1058.000000	1058.000000	
	1058.000000		
mean	75.112543	16.443195	
4.987713			
std	14.496912	6.662057	
2.831910			
min	50.030000	5.000000	
1.000000			
25%	62.595000	11.200000	
2.000000			
50%	74.940000	15.800000	
5.000000			
75%	87.882500	21.600000	
7.000000			
max	99.960000	30.000000	
10.000000			

Cluster 2:

	Midterm Score	Final Score	Assignments Score (Averaged) \
count	1064.000000	1064.000000	1064.000000
mean	69.218882	84.760254	74.614870
std	15.918761	9.577428	13.490179
min	40.020000	56.820000	50.000000
25%	56.607500	78.210000	64.002500
50%	67.980000	85.755000	74.828117
75%	82.115000	92.390000	84.745000
max	99.880000	99.980000	99.780000

	Quizzes Score (Averaged)	Study Hours per Week	Stress Level
(1-10)			
count	1064.000000	1064.000000	
	1064.000000		
mean	77.387284	24.202256	
5.535714			
std	14.446141	3.852310	
2.857981			
min	50.190000	13.400000	
1.000000			
25%	65.560000	21.300000	
3.000000			
50%	78.415000	24.700000	
6.000000			

75%	90.130000	27.400000	
8.000000			
max	99.940000	30.000000	
10.000000			
Cluster 3:			
	Midterm Score	Final Score	Assignments Score (Averaged) \
count	1062.000000	1062.000000	1062.000000
mean	84.840377	54.583004	74.637426
std	9.587196	9.023724	13.728215
min	61.530000	40.000000	50.010000
25%	77.452500	47.107500	62.910000
50%	85.580000	53.810000	74.828117
75%	93.227500	60.782500	85.842500
max	99.970000	81.440000	99.920000

	Quizzes Score (Averaged)	Study Hours per Week	Stress Level
(1-10)			
count	1062.000000	1062.000000	
1062.000000			
mean	75.923399		19.177589
5.794727			
std	14.476284		6.773679
2.865711			
min	50.160000		5.000000
1.000000			
25%	63.620000		13.900000
3.000000			
50%	76.630000		19.500000
6.000000			
75%	88.385000		25.000000
8.000000			
max	99.960000		29.900000
10.000000			